

# ANIA

## Assessment Report

A critical analysis of Portugal's  
**National Artificial Intelligence Agenda**  
from the social-cohesion perspective

June 2026

António Cruz  
Bruno Santos  
Diana Matos  
Pedro Miranda





## Table of Contents

|                                                                      |    |
|----------------------------------------------------------------------|----|
| 1. Introduction .....                                                | 3  |
| 1.1 Why AI and social cohesion in 2026 .....                         | 3  |
| 1.2 Framework .....                                                  | 3  |
| 1.3 The subject: ANIA in the Portuguese policy stack.....            | 3  |
| 1.4 Scope and approach .....                                         | 4  |
| 2. ANIA assessed through the social-cohesion lens .....              | 5  |
| 2.1 Digital inclusion and AI literacy .....                          | 5  |
| 2.2 Algorithmic fairness in public services .....                    | 5  |
| 2.3 Labour-market inequality.....                                    | 6  |
| 2.4 Cultural and linguistic cohesion .....                           | 7  |
| 2.5 Regional and territorial cohesion .....                          | 7  |
| 2.6 The structural-gap finding .....                                 | 7  |
| 3. Alignment with EU strategic documents.....                        | 8  |
| 3.1 The EU instruments and ANIA's alignment claim .....              | 8  |
| 3.2 Where the alignment is operationally thin .....                  | 8  |
| 4. Comparison with eight national AI strategies .....                | 9  |
| 4.1 Why these eight, and how to read them .....                      | 9  |
| 4.2 APAC: small-state pragmatism and population-scale ambition ..... | 9  |
| 4.3 Europe: humanism as method, twin transition as horizon .....     | 10 |
| 4.4 Americas: economic opportunity and inclusion-as-purpose .....    | 11 |
| 4.5 Africa: distributed institutions and measured ambition .....     | 11 |
| 4.6 Overall comparison with ANIA.....                                | 12 |
| 5. Critical assessment and targeted recommendations .....            | 13 |
| 5.1 Synthesis of findings .....                                      | 13 |
| 5.2 Targeted recommendations .....                                   | 13 |
| 5.3 Sequencing .....                                                 | 14 |
| 6. Conclusion .....                                                  | 15 |
| References .....                                                     | 16 |

# 1. Introduction

## 1.1 Why AI and social cohesion in 2026

Two Portuguese baselines make the social-cohesion question concrete before AI deployment enters it. The first is the grey divide. Forty per cent of Portuguese aged sixty-five to seventy-four do not use the Internet (Castells, 2024), so any AI-mediated citizen service starts from a population in which a substantial cohort is digitally absent. Whether a public service that migrates to AI-powered chatbots reaches that cohort, or excludes it further, is decided by institutional choice, not by the technology itself.

The second baseline is discrimination where citizens meet the state. Eighty-four per cent of Portuguese Roma (members of the Romani ethnic minority) report perceived ethnic profiling in police stops (FRA, 2022), and any predictive policing or eligibility-screening tool trained on that police data inherits the bias. Whether such a tool is deployed, and under what oversight, is also an institutional choice.

National AI strategies are the principal instrument coordinating those choices: who is trained, which public services adopt AI under what safeguards, which languages are preserved in foundation models, which regions are reached, and which groups are heard. This report assesses Portugal's National AI Agenda, ANIA, through that distributional lens.

## 1.2 Framework

The cohesion question for AI deployment is operationalised here as the quality of relations between citizens, and between citizens and the institutions that govern them, following Schirch (2023). That quality is too broad to assess directly, so the report decomposes it into five sub-dimensions: digital inclusion and AI literacy; algorithmic fairness in public services; labour-market inequality; cultural and linguistic cohesion; regional and territorial cohesion.

The decomposition is operational rather than canonical. Each sub-dimension maps to identifiable provisions in the European Union's AI strategy and to identifiable initiatives inside ANIA, which makes the assessment a structured comparison, not an interpretive essay.

## 1.3 The subject: ANIA in the Portuguese policy stack

Portugal published its National AI Agenda in January 2026 (Resolution n.º 2/2026; República Portuguesa, 2026a). ANIA sits as Action 20 of the broader Estratégia Digital Nacional (EDN) Action Plan 2026 to 2027 (República Portuguesa, 2026b), the umbrella digital-transformation strategy of which ANIA is the AI-specific track.

ANIA articulates four Areas of Action:

- Area I (infrastructure and data)
- Area II (innovation and adoption)
- Area III (talent and skills)
- Area IV (responsibility and ethics)

The four Areas are operationalized through thirty-two named initiatives anchored on six guiding principles; each initiative is numbered within its Area, so Initiative III.8 is the eighth initiative of Area III. The EDN itself contains a separate sequence of twenty named EDN Actions referenced occasionally below, for example EDN Action 2.5 on partnerships with Portuguese-official-language countries or EDN Action 13.4 on the Continent-Azores-Madeira and Azores submarine-cable rings.

The Agency for the Technological Reform of the State (ARTE), established by Decree-Law n.º 96/2025, complements ANIA by hosting the future AI Centre of Excellence for Public Administration (República Portuguesa, 2025a). ANIA presents itself as aligned with three European-Union instruments: the AI Act (Regulation (EU) 2024/1689), the AI Continent Action Plan, and the Apply AI Strategy.

## 1.4 Scope and approach

This report is the social-cohesion deliverable of the *Inteligência Artificial na Sociedade* course at ISCTE. The other three perspectives the course covers, work, democracy, and ethics and human development, enter the analysis only where they illuminate the cohesion question.

The comparative work was built from a wider corpus of twenty-three country analyses and then narrowed deliberately to two comparators per world region, capped at eight. Asia-Pacific is represented by Australia together with New Zealand and by India; Europe by France and Finland; the Americas by the United States and Brazil; Africa by South Africa and Egypt.

The four-region grid forces the comparison across very different development trajectories, demographic profiles, and political traditions instead of clustering inside the European and OECD frontier where Portugal already sits. The eight strategies are read as a map of the design space ANIA could occupy, not as a ranking. A SWOT and a PESTLE annex consolidate the comparative findings.

## 2. ANIA assessed through the social-cohesion lens

ANIA's architecture comprises four Areas of Action and thirty-two named initiatives anchored on six guiding principles, and the strategy positions itself as aligned with the AI Continent Action Plan and the Apply AI Strategy. The five sub-dimensions of cohesion introduced in §1.2 supply the lens for the assessment that follows.

### 2.1 Digital inclusion and AI literacy

The strongest move ANIA makes on digital inclusion is the cross-ministerial mandate for citizen AI literacy under Area III. Initiative III.8 (AI Literacy for Citizens) places the literacy action in FCT with named co-leadership from ANACOM, CNCS, IPQ, and ARTE, reinforced by Initiative III.6 (National AI Week), Initiative III.7 (the AI Generation youth campaign), and Initiative III.2 (the National Smart Skills Framework). The mandate exists; what it lacks is a quantified per-cohort commitment. The Cedefop benchmark of sixty-one per cent of adults with basic digital skills, including fifteen per cent of older adults, implies several hundred thousand Portuguese reached annually if the grey-divide baseline noted in §1.1 is to close. ANIA names no such target.

The chosen comparators provide three concrete templates of what a quantified commitment looks like. India's IndiaAI FutureSkills enrolls eight hundred and sixty thousand and pairs the platform with Responsible AI for Youth reaching more than fifty thousand secondary-school students. South Africa commits to thirty per cent whole-population workforce literacy plus twenty thousand specialists and five thousand experts. Egypt names a thirty-six per cent public-access KPI alongside a twenty-six per cent marginal-population target, the only marginal-population quantification in the comparator set. ANIA's silence on cohort coverage leaves the question of who gets reached open to discretion downstream of the strategy.

### 2.2 Algorithmic fairness in public services

The regulatory perimeter ANIA builds around AI Act compliance is coherent on the page. Initiative IV.4 (EU AI Act transposition) anchors the national application; Initiative IV.6 (the EU AI Act implementation guide) is produced under joint leadership of ANACOM, ARTE, IPQ, and ANI; Initiative IV.3 (the Centre for Responsible AI) provides the institutional locus; and the AI Centre of Excellence inside Public Administration is placed by Initiative II.11.

The structural weakness is that the perimeter does not name the operational instruments that comparator jurisdictions use to make algorithmic fairness enforceable in practice.

Australia commits to a recurring multi-watchdog AI Review Committee that meets every six weeks and aggregates the Information Commissioner, the Privacy Commissioner, the Ombudsman, and rotating sectoral regulators. Brazil plans a National Centre for Algorithmic Transparency under Action 51 of its AI Plan, and the broader Brazilian governance practice, which involved over one hundred institutions in the design of the strategy, illustrates what an AI Ethics Expert Group with structured civil-society representation looks like in practice. None of these instruments is named in ANIA's responsibility pillar.

The standing reminder of why operational machinery matters is the Dutch *toeslagenaffaire*, in which twenty-six thousand mostly immigrant-background families were wrongly accused of fraud by a biased algorithm in welfare administration (FRA, 2022). Abstract principles did not prevent that outcome, and abstract principles inside ANIA will not prevent its Portuguese equivalent. A deeper concern follows from Schirch (2023): the responsibility pillar's orientation is compliance-and-content-moderation, and that orientation is structurally insufficient because the harms AI deployment produces emerge from design choices made upstream of moderation.

The recommendations developed in §5 push toward upstream instruments: a multi-watchdog Review Committee, an Ethics Expert Group with civil-society participation, and an open-source algorithmic-fairness toolkit ANIA could adopt instead of building from scratch.

## 2.3 Labour-market inequality

On the workforce side, ANIA's commitments cluster around Doutor AP (Initiative III.1, the accelerated AI training programme for advanced civil servants), a set of skills initiatives in Area III, and research funding for work-impact and algorithm-auditing studies under Area IV. The strategy commits effort, but it leaves several openings that the comparator strategies close.

Coverage is the first. Doutor AP trains an advanced cohort, where Australia instead mandates foundational AI training for the entire Australian Public Service within twelve months of adoption. Beyond coverage, ANIA lacks an institutional counterpart to the trade-union embedding that France and Finland both build into governance: France through the social-dialogue mandate inside its Humanism pillar, Finland through trade-union representation in the strategic-level roundtable proposed by the AI 4.0 Programme. ANIA names no comparable structure for organised labour in the governance of AI's workforce effects.

Monitoring is the deeper absence. The Agenda commits to no system for tracking how AI's labour-market effects fall across gender, age, disability, or migrant origin. Brazil binds itself to a fifty per cent women-researcher target as a hard constraint on its talent pipeline. India mainstreams gender through Jan Dhan (the *Pradhan Mantri Jan Dhan Yojana* national financial-inclusion programme) while naming four hundred and ninety million informal workers as an explicit policy population. Without that monitoring, the strategy cannot tell whether AI is closing or widening the inequalities it was nominally meant to address.

This is not only a measurement problem but a conceptual one. Therborn (2009) distinguishes three kinds of inequality (vital, existential, material) and four mechanisms that produce them: distantiation, exclusion, hierarchy, exploitation. AI deployment carries material-inequality risk through distantiation, where front-runners pull ahead while the rest fall behind without explicit barriers, and existential-inequality risk through algorithmic mis-recognition, of which the Roma profiling pattern noted in §1.1 is a documented Portuguese instance. Naming those mechanisms is the precondition for monitoring them; ANIA's silence on either leaves the cohesion consequences un-named.

## 2.4 Cultural and linguistic cohesion

ANIA's cultural-and-linguistic move is concentrated in Initiative II.7 (AMALIA, the sovereign European-Portuguese language model), paired with Initiative I.3 (sectoral data spaces in critical areas). The pair aligns with the European Union's ALT-EDIC and is complemented by EDN Action 2.5, which commits to partnerships with Portuguese-official-language countries. The strength is a named sovereign model with explicit institutional partners.

Three structural gaps remain. AMALIA is not paired with a sectoral data-sets initiative analogous to Egypt's Initiative 7, which builds high-quality Arabic data sets specifically as the substrate the model needs in order to serve Arabic speakers across domains. ANIA does not commit AMALIA to open-source or open-weight distribution, and that choice will shape whether Portuguese institutions can audit, fine-tune, and embed the model in their own systems. And there is no Atlantic-Lusophone cultural-heritage data programme anchoring AMALIA in the broader Portuguese-speaking world, a gap that becomes structurally significant once §4 introduces Brazil's PT-BR sovereign model as the natural bilateral coordination partner.

## 2.5 Regional and territorial cohesion

Territorial cohesion is the dimension where the ANIA-internal architecture is weakest. The compute and data initiatives in Area I are nationally anchored without explicit regional distribution: the Advanced Computing AI Factory with Deucalion expansion (I.1), the Gigafactory (I.2), and the National Data Centre Plan (I.4).

The territorial dimension is carried by the broader EDN instead of by ANIA itself, through Mobile Citizen Service Points under ARTE leadership (Action 11.4) and the CAM Ring and Azores Ring submarine cables (Action 13.4). The gap inside ANIA is the absence of any AI initiative framed explicitly around territorial cohesion. South Africa's AIISA model is the most direct contrast in the chosen set: eleven AI Hubs distributed across universities outside the largest metropolitan centre, designed precisely so that AI capacity does not concentrate in a single national capital.

A National AI Hub Network sketched on that model, with named host universities including the Açores and Madeira, is one of the recommendations §5 returns to.

## 2.6 The structural-gap finding

Across the five sub-dimensions, ANIA addresses cohesion at the level of intent, through guiding principles and EU AI Act compliance, but does not establish a structural cohesion pillar. Among the chosen eight, four treat cohesion as a central or near-central operational concern (the Australia and New Zealand pair, France, Brazil, and India), and only the United States treats it as a derivative of competitiveness in the way ANIA does. The Portuguese gap is not unique among European peers, but the structural-cohesion-pillar move is what distinguishes the strategies that have made cohesion operational from those that have not.

## 3. Alignment with EU strategic documents

### 3.1 The EU instruments and ANIA's alignment claim

ANIA presents itself as aligned with two 2025 European Union instruments. The alignment is strong on rhetoric and thinner on operational specificity, in ways §3.2 returns to. The AI Continent Action Plan (European Commission, 2025a) articulates five pillars covering compute infrastructure, data, AI algorithms and adoption, skills, and regulatory simplification.

The Apply AI Strategy (European Commission, 2025b) structures EU AI policy around eleven sectoral flagships, three cross-cutting challenges (small and medium enterprises, the AI-ready workforce, and frontier AI), and a single governance mechanism anchored by an AI Observatory.

Neither instrument treats social cohesion as a standalone pillar; both treat it as a cross-cutting evaluation criterion attached to skills, public-sector adoption, healthcare deployment, accessibility, and workforce protection. ANIA adopts that same cross-cutting framing, which is the principal basis on which the alignment claim holds.

### 3.2 Where the alignment is operationally thin

The alignment weakens in four related places, all of which map onto the operationalisation gap §2 has already surfaced inside ANIA.

The level of commitment is the first. ANIA's principles sit at the level of intent, where the EU documents commit to named monitoring instruments such as the AI Observatory under the AI Office, the AI Literacy repository under Article 4 of the AI Act, and the AI Skills Academy inclusion reporting.

Time discipline is the second. The AI Continent Action Plan dates its actions to specific quarters of 2025, while ANIA initiatives carry semester-level start dates without semester-level deliverable or cohort-coverage targets, which leaves the strategy without an external clock against which to track its own delivery.

Accessibility and persons with disabilities are the third. They sit explicitly in the EU framing and implicitly in ANIA's guiding principles, a difference that becomes operationally significant once comparator strategies turn the principle into named instruments.

And the substrate problem already noted in §2.4 reappears here: ANIA does not pair AMALIA with a sectoral data-sets commitment of the kind Egypt provides for Arabic, leaving the data foundation of Portuguese-language AI sovereignty under-articulated.

These four gaps are not failures of intent but failures of operationalisation, addressed in §5 by extending existing initiatives in Areas III and IV and EDN Action 2.5 instead of by adding new institutional layers.



## 4. Comparison with eight national AI strategies

### 4.1 Why these eight, and how to read them

The comparison began from a wider corpus of twenty-three country analyses and was narrowed deliberately to two comparators per world region, capped at eight. Australia together with New Zealand and India represent the Asia-Pacific; France and Finland represent Europe; the United States and Brazil represent the Americas; South Africa and Egypt represent Africa. The four-region grid is the analytical move that matters: it forces the comparison across very different development trajectories, demographic profiles, and political traditions instead of clustering inside the European and OECD frontier where Portugal already sits. The eight strategies are read as a map of the design space ANIA could occupy, not as a ranking.

Within each pair, the two comparators were chosen to stage a deliberate analytical contrast, not to represent the region neutrally.

Asia-Pacific juxtaposes the small-advanced-economy adoption strategy (Australia together with New Zealand) against the population-scale ambition of India, the two extremes of the design space at opposite demographic and developmental scales. Europe pairs France's humanism-first organisation of the strategy with Finland's twin-transition framing, isolating cohesion-as-method against cohesion-as-horizon inside the same EU-AI-Act regulatory perimeter. The Americas juxtapose the United States' competitiveness-derivative framing with Brazil's inclusion-as-operational-purpose framing, the sharpest ideological contrast in the chosen set. Africa pairs South Africa's distributed-institutional model with Egypt's centralised measurement discipline, the two distinct continental positions on whether cohesion is built through architecture or through metrics.

Every pair therefore answers a different version of the same question: how does the strategy organise itself when social cohesion is the lens? The four regional sections below argue the reasoning behind each pair, and §4.6 returns to what the comparison says about ANIA.

### 4.2 APAC: small-state pragmatism and population-scale ambition

The Asia-Pacific (APAC) pair sets two very different strategic logics against each other. Australia and New Zealand are treated together because they share a common-law tradition, comparable demographic scale, and explicit indigenous-rights anchoring. Together they represent the small-advanced-economy adoption strategy, and neither attempts to compete with the United States or China on foundational model development.

Australia frames its position through “shared prosperity”: the December 2025 National AI Plan places benefit distribution on equal footing with productivity, supports First Nations data sovereignty, and commits to union involvement in workforce transitions.

New Zealand pushes the indigenous-rights dimension further, embedding the Treaty of Waitangi into its July 2025 strategy and treating *mātauranga Māori* as cultural property protected from extractive training; Te Hiku Media's *te reo Māori* speech-recognition model

is cited as the national exemplar of that posture. The regulatory choice in both is deliberately light-touch and OECD-Principles-aligned, in explicit rejection of an AI-Act equivalent.

India sits at the opposite end of the same region. The IndiaAI Mission, approved by the Union Cabinet in March 2024 with an outlay of about ₹10,371 crore (roughly US\$1.25 billion), is the most ambitious cohesion programme in the chosen set, framed throughout by the #AIForAll vision and anchored to a demographic dividend running into the 2050s.

Three Indian moves have no parallel anywhere else in the set. BHASHINI provides government services in more than thirty-six Indian languages and BharatGen is a sovereign multimodal foundation model covering twenty-two of them. SabhaSaar and eGramSwaraj take AI down to the village-assembly level, with over two hundred and fifty thousand gram panchayats onboarded in a single fiscal year. And Digital ShramSetu makes India the only country in the set that explicitly targets the informal economy with AI-supported livelihood services.

For ANIA, the Asia-Pacific pair surfaces three lessons Portugal underweights. The most immediate is the cultural-and-linguistic anchor that New Zealand builds through Māori embedding and India through regional-language coverage, which Portugal could extend to its own variants and to immigrant-resident languages. The deeper lesson is the scale at which linguistic-inclusion infrastructure can be built, as India operationalises through BHASHINI and BharatGen.

The structural lesson is the explicit deployment of AI to local levels of governance and to informal-economy workers, two policy spaces ANIA's text does not name.

### 4.3 Europe: humanism as method, twin transition as horizon

France is the European strategy in the set that most explicitly organises its vision around social cohesion. The French Artificial Intelligence Commission report of March 2024 structures itself around three pillars: Humanism, Sovereignty, and Responsibility. The document argues that without social adoption of AI there is a risk of mass rejection, drawing a parallel with past technological revolutions that only delivered productivity gains when work organisations were restructured around them. Social cohesion is therefore positioned as a precondition for economic success, not as a derivative effect, which stands in direct contrast to ANIA's framing.

The Humanism pillar is operationalised concretely. Social dialogue and co-construction are presented as conditions for AI to improve, not degrade, working conditions, with the document going as far as proposing that AI itself be deployed as a tool to support that dialogue. The *AI Cafés* are physical spaces for public experimentation and debate on AI for ordinary citizens, part of a broader national awareness plan that includes a digital information tool and a competition for positive AI use cases. The training plan is differentiated by group, with specific measures for students, employees, civil servants, and retirees, recognising that AI literacy cannot be treated as a single homogeneous audience. The document also estimates that France's annual economic growth could

double through AI adoption but that this effect would be temporary, making investment in cohesion and training a structural necessity instead of an ethical option.

Finland integrates AI into a broader Twin Transition in which digital and green transformations are pursued as a single integrated objective, not as parallel tracks. The Artificial Intelligence 4.0 Programme proposes a strategic-level roundtable with explicit trade-union representation as the governance mechanism, in place of a permanent agency. Finland is also the only comparator in the chosen set to use quantitative international-ranking targets as its measurable performance horizon, which gives the strategy a tracking discipline that ANIA currently lacks.

The European pair offers two transferable moves for ANIA: France’s deliberation-and-dialogue model folded into Area III alongside existing reskilling commitments, and Finland’s measurable-performance discipline tracked against named external indices.

#### 4.4 Americas: economic opportunity and inclusion-as-purpose

The Americas pair illustrates the sharpest ideological contrast in the chosen set. The United States’ *America’s AI Action Plan*, published in July 2025, treats AI primarily as an instrument of national competitiveness. The plan devotes attention to workers through a “worker-first” reskilling agenda, tax incentives for AI training, and a Labour Market Analysis Hub, but the framing is consistently economic. References to algorithmic bias and to diversity, equity, and inclusion were actively removed from the NIST AI Risk Management Framework and from federal procurement guidance for large language models as part of the plan’s review. The United States model demonstrates one extreme of the design space: cohesion pursued through opportunity creation, not through protection.

Brazil sits at the opposite pole. The 2024 to 2028 *AI for the Good of All* plan is the most explicitly inclusion-oriented strategy in the corpus. It commits the largest absolute investment in the set, approximately R\$23 billion (roughly €4 billion), and reaches a population scale unmatched elsewhere. The *Acredita no Primeiro Passo - IA* initiative, built on the CadÚnico registry of low-income households, is designed to reach up to ninety-seven million people. A fifty per cent women-researcher target is binding. Training programmes target vulnerable populations explicitly. Multi-sectoral governance involved over one hundred institutions in design, and a public *Plataforma de Participação Social* anchors continuing consultation. Brazil also commits to a state-funded PT-BR language model. The distinctive Brazilian move is to name inclusion as the operational purpose of the strategy rather than as a derivative of competitiveness.

ANIA currently sits closer to the United States framing than to the Brazilian one, despite Portugal’s social-protection tradition. Brazil’s PT-BR commitment opens a bilateral coordination opportunity around AMALIA that has no European-side equivalent in the chosen set.

#### 4.5 Africa: distributed institutions and measured ambition

South Africa’s AI Planning (October 2023), anchored to the National Development Plan 2030, treats jobs displacement explicitly and commits to thirty per cent whole-population



workforce literacy alongside twenty thousand specialists and five thousand experts. The South African signature move is institutional: the AI Institute of South Africa operates through eleven AI Hubs in universities outside the largest metropolitan centre, a deliberately distributed model rather than a concentrated one. Continental cooperation under the African Continental Free Trade Area reframes cohesion as a regional, not a national, concern. The hub-network design contrasts directly with the narrower single-centre model Portugal currently runs through the Centre for Responsible AI.

Egypt's 2025 to 2030 strategy is the most measurement-disciplined in the corpus. A bespoke eighty-nine-indicator Capacity Assessment Framework produces a single composite national score and tracks an explicit global-ranking trajectory. Two paired initiatives anchor linguistic cohesion: Arabic-first foundation models, and high-quality sectoral Arabic data sets supplying the substrate the model needs to serve Arabic speakers. The strategy names a twenty-six per cent marginal-population KPI as well, the only comparator in the chosen set that quantifies a vulnerable-population commitment at the strategy level. The Egyptian framework is portable: a Capacity Assessment instrument and a paired model-and-data approach would translate cleanly to AMALIA and to a Portuguese-language data substrate.

## 4.6 Overall comparison with ANIA

Four anchor points emerge that \$5 turns into ANIA-shaped recommendations: cohesion as a stated purpose rather than as a derivative effect (the Australia and New Zealand pair, Brazil, France, India); operational specificity in named instruments (AIISA, BHASHINI and BharatGen, the Capacity Assessment Framework, CadÚnico programmes, AI Cafés); substrate-level linguistic specificity (India, Egypt, New Zealand); and embedded participation in design and governance (Brazil, France, Australia, Finland). The chosen comparators turn each from principle into operational commitment, where ANIA so far does not.

## 5. Critical assessment and targeted recommendations

### 5.1 Synthesis of findings

ANIA is coherent and EU-anchored, but its treatment of social cohesion is structurally weaker than the leading comparators: cohesion sits at the level of guiding principles, not at the operational core of the strategy. Among the chosen eight, only the United States treats cohesion as a derivative of competitiveness in the way ANIA does. Brazil, France, India, and the Australia and New Zealand pair treat it as central or near-central, and the most operationally specific instruments in the corpus come from outside the European core, particularly from Brazil, India, South Africa, and Egypt. The nine recommendations below extend existing ANIA initiatives and EDN Actions without requiring a Resolution amendment.

### 5.2 Targeted recommendations

The nine recommendations organise around the four anchor points surfaced in §4.6. Each is presented as a named instrument followed by its rationale; together they extend existing ANIA initiatives and EDN Actions without requiring a Resolution amendment.

#### **Anchor 1: cohesion as operational purpose.**

- **Society and Inclusion sub-pillar within Area IV.** Add a named cross-cutting sub-pillar inside Area IV (Responsibility and Ethics) carrying its own indicators and reporting line, in the spirit of Brazil's *AI for the Good of All* framing. The move repositions social cohesion from the guiding-principles layer into the strategy's architecture, where it can be tracked and audited. It needs no Resolution amendment and can be introduced through the H1 2027 mid-cycle review.

#### **Anchor 2: operational specificity.**

- **Tripartite Observatory** representing UGT, CGTP, CIP, and academia, modelled on France's social-dialogue mandate and Finland's strategic-level roundtable, extending Area IV's labour-impact research strand into a standing institution.
- **Multi-watchdog Review Committee** meeting every six weeks, aggregating CNPD, the Provedor de Justiça, CADA, and rotating sectoral regulators, adopting the Australian model under the AI Act transposition track.
- **National AI Hub Network** of seven to ten hubs at named host universities including the Açores and Madeira, paired with CPLP-Atlantic cooperation, importing the South African AIISA distributed model.
- **Egyptian-style Capacity Assessment Framework** producing a single composite national score with a target of EU top-five by 2030, paired with an ANIA Second Edition for late 2027, importing Egypt's measurement discipline as a new cross-cutting commitment.

**Anchor 3: linguistic and cultural specificity at the substrate level.**

- **Bilateral AMALIA coordination** with Brazil's PT-BR work, paired with a sectoral Portuguese data-sets initiative built on the Egyptian model-and-data approach. Closes the substrate gap identified in §2.4 and §3.2 by extending the AMALIA cluster across Areas I and II and EDN Action 2.5.
- **Open-source algorithmic-fairness toolkit** adoptable inside the Centre for Responsible AI. Closes the fairness gap identified in §2.2; Singapore's *AI Verify* is the available reference instrument and is retained here as a deliberate scope-exception because no chosen comparator supplies an equivalent open-source toolkit.

**Anchor 4: participation in design and governance.**

- **Participatory drafting methodology** with regional Town Halls and named civil-society convenors, paired with a Portuguese AI Cohesion Monitor as the ongoing tracking instrument. Draws on Brazil's broad multi-stakeholder governance practice and on France's social-dialogue mandate.
- **Gender and disability quotas** in the Talent-pillar workforce targets (fifty per cent women and named representation for persons with disabilities), paired with an AI Ethics Expert Group convened with civil-society participation. Draws on Brazil's women-researcher target and India's Jan Dhan gender mainstreaming.

## 5.3 Sequencing

The sequencing follows three waves within the 2026 to 2030 ANIA horizon. The first wave (H2 2026) launches the Review Committee, fairness toolkit, Talent-pillar quotas, and AMALIA coordination MoU negotiations. The second wave (H1 2027) brings the Society and Inclusion sub-pillar, the Ethics Expert Group, and the Observatory's first annual report. Deeper structural reforms, covering the Hub Network, participatory drafting, Capacity Framework, and ANIA Second Edition, extend through 2027 and 2028.

## 6. Conclusion

Whether AI is deployed to support social cohesion or to erode it depends on institutional choices, not on the technology itself. Treating cohesion as the quality of horizontal and vertical relations between citizens and between citizens and institutions makes the question tractable, because every initiative in a national AI strategy can then be assessed against whether it expands or contracts citizen agency, intergroup trust, or institutional legitimacy. The five-sub-dimension spine applied in this report converts that abstract question into a structured comparison.

Applied to ANIA, the comparison yields a consistent finding: the Agenda is regulatorily coherent but operationally weak on cohesion. The strategy folds cohesion into guiding principles, where the most operationally specific comparator strategies, those of Brazil, India, South Africa, and Egypt, build named instruments instead: distributed AI hub networks, civil-society-convened ethics expert groups, gender and disability quotas in workforce targets, paired model-and-data linguistic sovereignty, participatory drafting methodologies, and single-composite measurement frameworks.

The nine recommendations developed in §5 lift specific instruments from the eight chosen comparators and anchor each to an existing ANIA initiative or EDN Action. Implemented as a set, they position Portugal distinctively at the intersection of the European Union, the Lusophone world, and the African continent. The next ANIA iteration is the moment at which those choices get made.



## References

Castells, M. (2024). *Advanced introduction to digital society* (chs. 7–9). Edward Elgar.

Commonwealth of Australia (2025). *National AI Plan*. Department of Industry, Science and Resources (December 2025).

Egypt MCIT (2025). *National Artificial Intelligence Strategy, Second Edition (2025–2030)*. Ministry of Communications and Information Technology, National Council for Artificial Intelligence.

European Commission (2025a). *AI Continent Action Plan* (COM(2025) 165 final). <https://digital-strategy.ec.europa.eu/en/library/ai-continent-action-plan>

European Commission (2025b). *Apply AI Strategy* (COM(2025) 723 final). <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52025DC0723>

Executive Office of the President (2025). *America’s AI Action Plan*. Office of Science and Technology Policy, The White House (July 2025).

FRA (European Union Agency for Fundamental Rights) (2022). *Bias in algorithms: Artificial intelligence and discrimination*. Publications Office of the European Union.

Government of Finland (2022). *Artificial Intelligence 4.0 Programme: Finland as a Leader in Twin Transition* (Ministry of Economic Affairs and Employment, 2022:63).

Government of France (2024). *Stratégie nationale en intelligence artificielle: Humanisme, Souveraineté, Responsabilité*.

MCTI (Ministério da Ciência, Tecnologia e Inovação) (2024). *Plano Brasileiro de Inteligência Artificial (PBIA) 2024–2028: IA para o Bem de Todos / AI for the Good of All*. Brasília: MCTI / CGEE.

MeitY (2024). *IndiaAI Mission* (approved 7 March 2024), building on NITI Aayog (2018), *National Strategy for AI: #AIForAll*. Government of India.

MBIE (Ministry of Business, Innovation and Employment) (2025). *New Zealand’s Strategy for Artificial Intelligence: Investing with Confidence*. Wellington: MBIE (July 2025).

Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). *Official Journal of the European Union*, L 1689.

República Portuguesa (2025a). *Decreto-Lei n.º 96/2025, de 21 de agosto, criando a Agência para a Reforma Tecnológica do Estado (ARTE)*. *Diário da República*, 1.ª Série.

República Portuguesa (2026a). *Agenda Nacional para a Inteligência Artificial (ANIA)*. Aprovada pela Resolução do Conselho de Ministros n.º 2/2026. <https://digital.gov.pt/pt/documentos/agenda-nacional-de-ia>

República Portuguesa (2026b). *Estratégia Digital Nacional: Plano de Ação 2026–2027*. <https://digital.gov.pt/pt/documentos/estrategia-digital-nacional>

Republic of South Africa (2023). *AI Planning: Adoption of AI by Government*. Department of Communications and Digital Technologies (discussion draft, October 2023).

Schirch, L. (2023). The case for designing tech for social cohesion: The limits of content moderation and tech regulation. *Yale Journal of Law & the Humanities*, 34, 1–34.

Singapore Government (2023). *National AI Strategy 2.0 (NAIS 2.0)*; AI Verify Foundation. Smart Nation and Digital Government Office. Cited as scope-exception.

Therborn, G. (2009). The killing fields of inequality. *Soundings*, 42, 20–32.